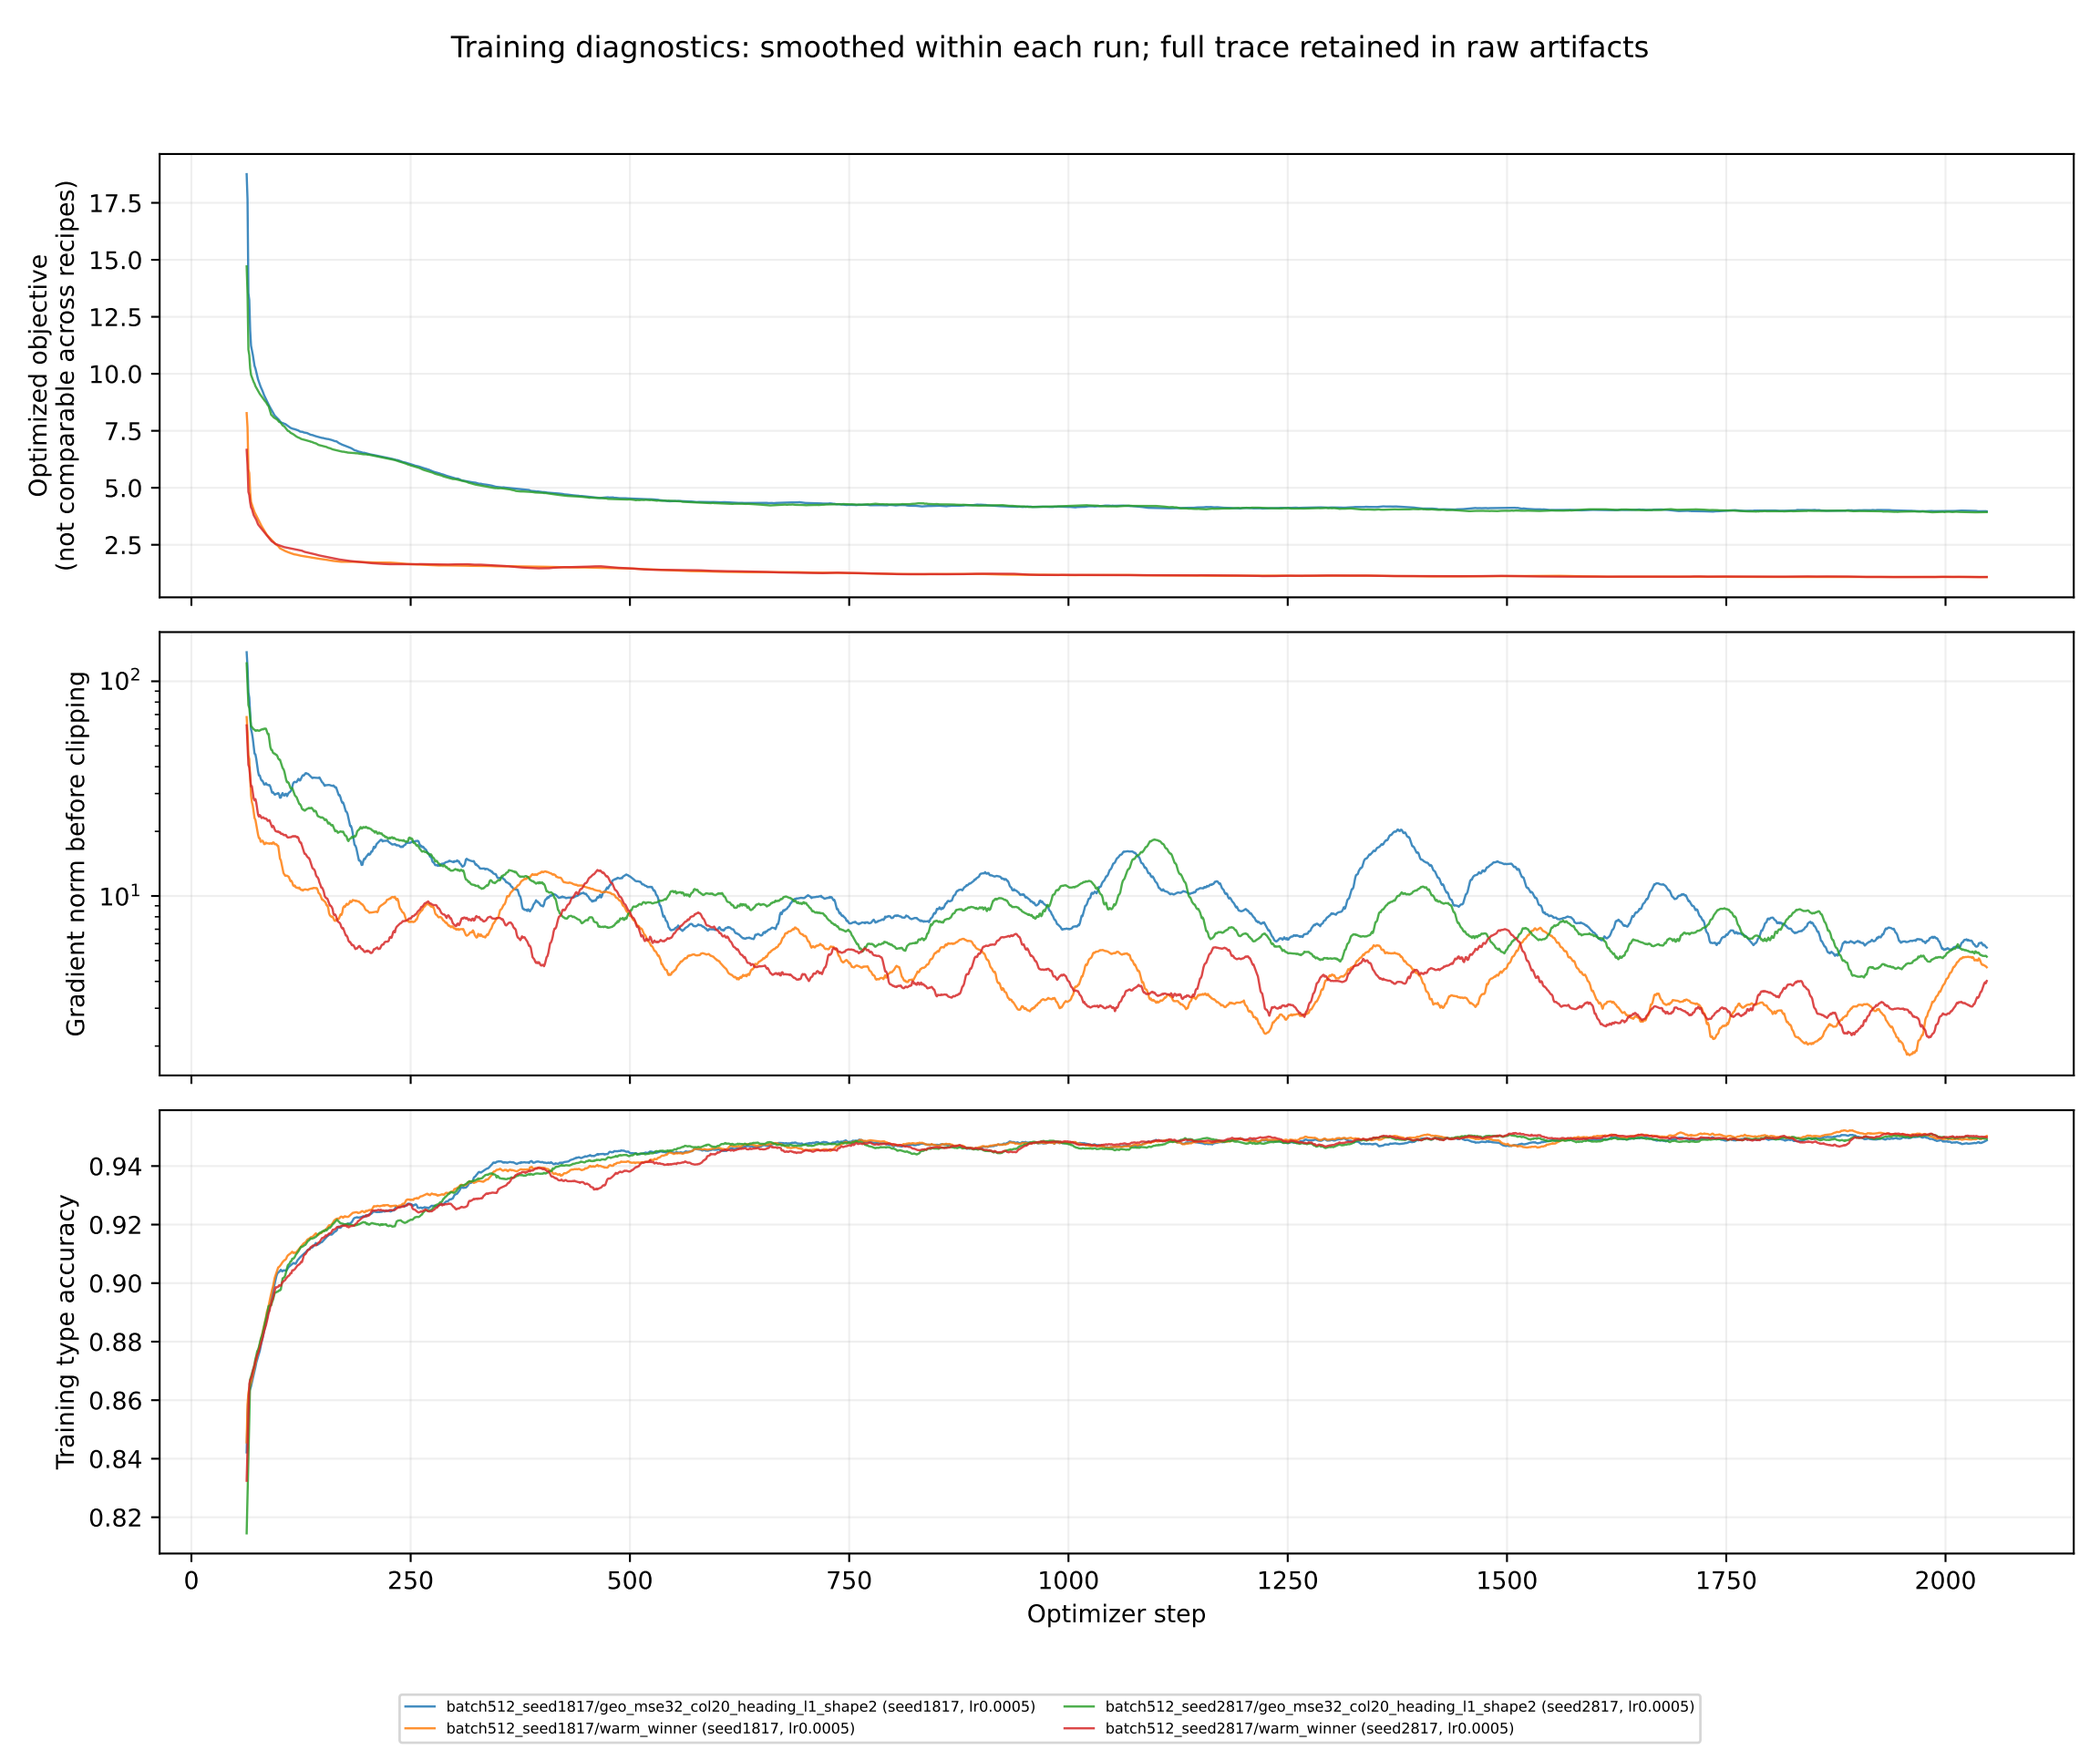

Training diagnostics: smoothed within each run; full trace retained in raw artifacts
Optimized objective
(not comparable across recipes)
17.5
15.0
12.5
10.0
7.5
5.0
2.5
Gradient norm before clipping
10^2
10^1
Training type accuracy
0.94
0.92
0.90
0.88
0.86
0.84
0.82
0
250
500
750
1000
1250
1500
1750
2000
Optimizer step
batch512_seed1817/geo_mse32_col20_heading_l1_shape2 (seed1817, lr0.0005)
batch512_seed1817/warm_winner (seed1817, lr0.0005)
batch512_seed2817/geo_mse32_col20_heading_l1_shape2 (seed2817, lr0.0005)
batch512_seed2817/warm_winner (seed2817, lr0.0005)